

Cython for HPC

Niall Moran



July, 2016

Overview

- ▶ Motivation
- ▶ Why (not) Python?
- ▶ Cython
- ▶ Examples



Hall Effect

- ▶ Edwin Hall (1879)
- ▶ Magnetic field induces Hall current





Hall Effect

- ▶ Edwin Hall (1879)
- ▶ Magnetic field induces Hall current





(Integer) Quantum Hall Effect

- ▶ von Klitzing (1980)
- ▶ Lower temperature, plateaus appear
- ▶ Quantization of conductance





(Integer) Quantum Hall Effect

- ▶ von Klitzing (1980)
- ▶ Lower temperature, plateaus appear
- ▶ Quantization of conductance





Fractional Quantum Hall Effect

- ▶ Tsui, Stormer and Gossard (1982)
- ▶ Plateaus at fractional filling
- ▶ Anyonic excitations



Fractional Quantum Hall Effect

- ▶ Tsui, Stormer and Gossard (1982)
- ▶ Plateaus at fractional filling
- ▶ Anyonic excitations



Modelling


- ▶ Finite system N particles and N_ϕ flux with $\nu = \frac{N}{N_\phi}$
- ▶ Decompose into magnetic orbitals



Ways to fill available orbitals

Modelling


- ▶ Finite system N particles and N_ϕ flux with $\nu = \frac{N}{N_\phi}$
- ▶ Decompose into magnetic orbitals



Ways to fill available orbitals

Modelling

- ▶ Finite system N particles and N_ϕ flux with $\nu = \frac{N}{N_\phi}$
- ▶ Decompose into magnetic orbitals




Ways to fill available orbitals

$$\binom{N_\phi}{N} \approx f(\nu)^N$$

Modelling

- ▶ Finite system N particles and N_ϕ flux with $\nu = \frac{N}{N_\phi}$
- ▶ Decompose into magnetic orbitals




Ways to fill available orbitals

$$\binom{N_\phi}{N} \approx f(\nu)^N$$

Wave-function is complex vector of this dimension!

Computations

- ▶ Linear algebra
- ▶ Markov-Chain Monte Carlo
- ▶ Tensor network methods



6 particles with 18 flux

Diminishing returns



10-12 electrons

Diminishing returns



10-12 electrons



16-20 electrons

Why (not) use Python

Cons

- ▶ Slow
- ▶ GIL
- ▶ Resource usage

Why (not) use Python

Cons

- ▶ Slow
- ▶ GIL
- ▶ Resource usage

Pros

- ▶ Fast development, debugging
- ▶ Amount of packages available
- ▶ Plotting
- ▶ Good glue

Cython



- ▶ Superset of python
- ▶ Compiles to C code
- ▶ Best of both worlds
- ▶ Good for wrapping existing C code

Provides speed increase by

- ▶ Compiling
- ▶ Providing explicit types and functions
- ▶ Can release the GIL

Hello World

Create hello.pyx containing

```
print("Hello world!")
```

and setup.py containing

```
from distutils.core import setup
from Cython.Build import cythonize

setup(
    ext_modules = cythonize("hello.pyx")
)
```

then run on command line

```
$ python setup.py build_ext --inplace
```

Notebook

Load the cython extension

```
%load_ext Cython
```

Can now create cython blocks

```
%%cython  
print("Hello world!")
```

Annotation

```
%%cython -a  
print("Hello world!")
```

Demonstration

```
def mean_plain(A, N):  
    sum = 0.0  
    j = 0  
    while j < N:  
        sum += A[j]  
        j+=1  
  
    return sum / N
```

220ms for array of 1 000 000 elements

```
cimport numpy as np
cimport cython

@cython.boundscheck(False)
@cython.wraparound(False)
def mean_cython4(
    np.ndarray[np.float64_t, ndim=1, mode="c"] A,
    int N
):
    cdef double sum = 0.0
    cdef int j = 0

    while j < N:
        sum += A[j]
        j += 1

    return sum / N
```

1ms for array of 1 000 000 elements

Matrix multiplication benchmark

	80x80	1500x1500
	Units: MFLOPS	
Optimal layout		
Python	0.94	0.98
Cython	1.08	1.12
Added types	179	177
boundscheck/wraparound	770	692
mode="c"/mode="fortran"	981	787
BLAS ddot (ATLAS)	1282	911
Intel C	2560	1022
gfortran $A^T B$	1113	854
Intel Fortran $A^T B$	2833	1023
NumPy dot	3656	4757
Worst-case layout		
Python	0.94	0.97
boundscheck/wraparound	847	175
BLAS ddot (ATLAS)	910	183
gfortran AB^T	861	94
Intel Fortran AB^T	731	94

Conclusions and thank you

Summary

- ▶ Significant speedup
- ▶ Not always optimal
- ▶ Useful when cannot be vectorised (ODEs, MCMC)
- ▶ Can wrap standard C libraries
- ▶ Can disable the GIL

Further reading

- ▶ Cython website, (<http://cython.org/>)
- ▶ D. S. Seljebotn, “*Fast Numerical Computation with Cython*”, SciPy 2009.
(http://conference.scipy.org/proceedings/SciPy2009/paper_2/full_text.pdf)
- ▶ S. Behnel et. al., “*Cython: The Best of Both Worlds*”, IEEE 2011.
(http://folk.uio.no/dagss/cython_cise.pdf)